

GB



ology@github

Predicting Player Features of the English Premier League with scikit-Learn

2018-07-09 BY GENE

In [a previous installment](#), I harvested football stats for every player in the English Premier League and saved them as CSV files (separated by tabs actually).

In this installment, I use the same data but predict the player rating (a floating point number) and field position (a string) with the techniques of linear regression and k nearest neighbors, respectively.

First, we read-in the data:

```
import pandas as pd
```

```
summary    = pd.read_csv('~/.sandbox/dev/Games/Football/Summary-2015-16-proce  
offensive  = pd.read_csv('~/.sandbox/dev/Games/Football/Offensive-2015-16-pro  
passing    = pd.read_csv('~/.sandbox/dev/Games/Football/Passing-2015-16-proce  
defensive  = pd.read_csv('~/.sandbox/dev/Games/Football/Defensive-2015-16-pro
```

Fortunately, by inspecting, we can see that each of these files is in the same player order (by Rating):

```
>>> print summary.shape
```

```
(298, 13)
```

```
>>> summary.head()
```

	Name	Player	Apps	Mins	Goals	Assists	Yel
R							
1	Riyad Mahrez	Leicester, 25, AM(CLR)	36(1)	3058	17.0	11.0	1.0
2	Dimitri Payet	West Ham, 29, M(CLR)	29(1)	2573	9.0	12.0	3.0
3	Alexis Sánchez	Arsenal, 27, M(LR),FW	28(2)	2446	13.0	4.0	1.0
4	Mousa Dembélé	Tottenham, 29, M(CR)	27(2)	2273	3.0	1.0	3.0
5	Mesut Özil	Arsenal, 27, M(CLR)	35	3049	6.0	19.0	4.0

	Red	SpG	PS%	AerialsWon	MotM	Rating
R						
1	NaN	2.3	73.6	0.9	10.0	7.84
2	NaN	2.3	80.2	0.1	5.0	7.74
3	NaN	3.6	79.6	0.7	6.0	7.72
4	NaN	0.8	90.0	1.1	4.0	7.69
5	NaN	1.4	86.3	0.1	6.0	7.66

```
>>> offensive.shape
```

```
(298, 14)
```

```
>>> offensive.head()
```

	Name	Player	Apps	Mins	Goals	Assists	SpG
R							
1	Riyad Mahrez	Leicester, 25, AM(CLR)	36(1)	3058	17.0	11.0	2.3
2	Dimitri Payet	West Ham, 29, M(CLR)	29(1)	2573	9.0	12.0	2.3
3	Alexis Sánchez	Arsenal, 27, M(LR),FW	28(2)	2446	13.0	4.0	3.6
4	Mousa Dembélé	Tottenham, 29, M(CR)	27(2)	2273	3.0	1.0	0.8
5	Mesut Özil	Arsenal, 27, M(CLR)	35	3049	6.0	19.0	1.4

	KeyP	Drb	Fouled	Off	Disp	UnsTch	Rating
R							
1	1.8	3.5	2.2	0.1	2.4	2.1	7.84
2	4.0	2.2	1.3	0.1	1.7	1.8	7.74
3	2.1	3.4	2.2	0.4	3.3	2.4	7.72
4	1.0	2.9	1.1	0.1	2.2	0.8	7.69
5	4.2	1.3	1.2	0.4	1.7	2.0	7.66

Next, after reading-in the data, is to put it in one single dataframe with all the

columns and zeros for empty ("na") values:

```
data = summary

feature_cols = list(offensive)
data[feature_cols] = offensive[feature_cols]

feature_cols = list(passing)
data[feature_cols] = passing[feature_cols]

feature_cols = list(defensive)
data[feature_cols] = defensive[feature_cols]

data = data.fillna(0)
```

There are a couple columns embedded in the "Player" value for each record. This includes the team name, the player age and the field position. This code extracts those for use in our modelling:

```
i = 0
for player in data.Player:
    i = i + 1
    [team, age, position] = player.split(" ")
    data.set_value(i, 'team', team)
    data.set_value(i, 'age', int(age))
    data.set_value(i, 'position', position)
```

```
import re
```

```
i = 0
for position in data.position:
    i = i + 1
    posn = re.sub('\(.\+\)', '', position)
    data.set_value(i, 'posn', posn)
```

Now the data features are all these:

```
>>> list(data)
['Name', 'Player', 'Apps', 'Mins', 'Goals', 'Assists', 'Yel', 'Red', 'SpG',
```

Inspecting a single row:

```
>>> data.loc[[1]]
```

	Name	Player	Apps	Mins	Goals	Assists	Yel
R							
1	Riyad Mahrez	Leicester, 25, AM(CLR)	36(1)	3058	17.0	11.0	1.0

	Red	SpG	PS%	...	Inter	Fouls	Offsides	Clear	Blocks	OwnG	\
R				...							
1	0.0	2.3	73.6	...	1.0	0.5	0.0	0.4	0.1	0.0	

	team	age	position	posn
R				
1	Leicester	25.0	AM(CLR)	AM


```
[1 rows x 34 columns]
>>> data.loc[[1]]['Rating']
```

1	7.84
---	------

```
Name: Rating, dtype: float64
```


Ok. So far so good. How about we visualize the data to get a feel for what relationships might hold?

```
import matplotlib.pyplot as plt
import seaborn as sns

response_col = ['Rating']

def pair_features():
    sns.pairplot(data, x_vars=feature_cols, y_vars=response_col, kind='reg')
    plt.show()

feature_cols = ['Mins', 'Goals', 'Assists', 'Yel', 'Red']
pair_features()
feature_cols = ['SpG', 'PS%', 'AerialsWon', 'MotM', 'KeyP']
pair_features()
feature_cols = ['Drb', 'Fouled', 'Off', 'Disp', 'UnsTch']
pair_features()
feature_cols = ['AvgP', 'Crosses', 'LongB', 'ThrB', 'Tackles']
pair_features()
feature_cols = ['Inter', 'Fouls', 'Offsides', 'Clear', 'Blocks']
pair_features()
feature_cols = ['OwnG', 'age']
pair_features()
```



This shows us that most of the numeric features will contribute to our ML training, but in addition to the non-numeric features, we can remove a couple that will probably

not contribute to meaningful prediction. Here is the code to remove those column names and set the **X** and **y** elements to the proper values:

```
feature_cols = list(data)
for col in ['Name', 'Player', 'Apps', 'Rating', 'team', 'position', 'Red', 'Off', '
    feature_cols.remove(col)

X = data[feature_cols]
y = data[response_col]
```

Now we can perform the Rating prediction with linear regression:

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1)

print X_train.shape, y_train.shape
print X_test.shape, y_test.shape

from sklearn.linear_model import LinearRegression

estimator = LinearRegression()
estimator.fit(X_train, y_train)

X_test.iloc[[1]]          # R=235
datum = data.iloc[[234]]  # Bakary Sako - M
X = datum[feature_cols]
y = datum[response_col]
estimator.predict(X)      # array([[6.63099864]])
```

This is a pretty close prediction! The true response value is:

```
>>> X_test.iloc[[1]] # R=235, Name=Bakary Sako, posn=M
>>> data.iloc[[234]].Rating
R
235    6.61
```

Ok! What about predicting a non-numerical response? For this, we will use the k-nearest-neighbors measure:

```
from sklearn.neighbors import KNeighborsClassifier

feature_cols = list(data)
for col in ['Name', 'Player', 'Apps', 'team', 'position']:
    feature_cols.remove(col)

response_col = ['posn']


estimator = KNeighborsClassifier(n_neighbors=18)
estimator.fit(X_train, y_train)

datum = data.iloc[[163]] # Gary Cahill - D
X = datum[feature_cols]
estimator.predict(X)      # array(['D'], dtype=object)

datum = data.iloc[[51]] # Erik Lamela - AM
X = datum[feature_cols]
estimator.predict(X)      # array(['AM', 'FW'], dtype=object)
```

Ha! Decent results!

Here is the accuracy of this model plotted for increasing values of k:



FILED UNDER: DATA, SOFTWARE
TAGGED WITH: FOOTBALL, PYTHON, SCIKIT-LEARN, SOCCER

Epistemologist-at-large

^ Top